

# EPA Vehicle Emission Score Prediction using Machine Learning

---

## Group 4:

Sayantan Majumdar

Dawit Wolday Asfaw

Cong Shen

Divya Reddy Manku

Akhil Reddy Annreddy

# Introduction

---

- Greenhouse gases (GHG) are responsible for trapping heat, thereby making the planet warmer.
- In the US, the transportation sector generates the largest share of GHG emissions (29% of 2019 GHG emissions) from burning fossil fuels.
- The Environmental Protection Agency (EPA) emissions score (smog rating) reflects vehicle tailpipe emissions (CO<sub>2</sub>) that contribute to local and regional air pollution

# Problem Description

---

- Building a machine learning model to predict the EPA emission score from different vehicles operating in the US.
- Classification problem, emission score: 1 (worst) - 10 (best)
- Compare different machine learning models to assess their performance on the data sets
- Broad business objective: Enabling policymakers to address critical issues in the sustainable transportation industry

# Data Source and Collection

---

- Available on the EPA fuel economy portal [<https://www.fueleconomy.gov/>]
- Two data sets:
  - Vehicle data (83 variables, 44075 observations)
  - Emissions data (8 variables, 42442 observations)
- The EPA has generated annual reports, but there are no existing publicly available notebooks having detailed machine learning workflows

<https://www.fueleconomy.gov/feg/pdfs/guides/FEG2021.pdf>

# Data Source and Collection

---

Data Set	Feature Name	Details
emissions	score	EPA 1-10 smog rating for fuelType1 (target variable)
vehicle	fuelType1	For single fuel vehicles, this will be the only fuel. For dual fuel vehicles, this will be the conventional fuel.
vehicle	highway08	Highway MPG for fuelType1
vehicle	barrels08	Annual petroleum consumption in barrels for fuelType1
vehicle	year	Model year
vehicle	VClass	EPA vehicle size class
vehicle	phevBlended	If True, this vehicle operates on a blend of gasoline and electricity in charge depleting mode

# Potential Issues & Challenges

---



Data Curation: Merging the Vehicles and Emissions data sets



Feature Engineering



Model Complexity: Hyperparameter Tuning



Evaluation Metrics

# Research Questions

---

- *Which predictors are most suited for predicting the EPA emission score?*
  - Which machine learning model works best for this data set?
  - Which vehicle class has poor emission ratings?

# Data Management

---

- After merging and cleaning data (removing missing and incorrect values), we have 28556 non-redundant observations.
- Data Transformation: Aggregating fuelType1 variable

```
In [153]: ve_df.fuelType1.value_counts()
```

```
Out[153]: Regular Gasoline    15596  
Premium Gasoline    12175  
Diesel              365  
Electricity         254  
Midgrade Gasoline   136  
Natural Gas         30  
Name: fuelType1, dtype: int64
```



We can combine midgrade and regular gasolines as 'Regular Gasoline'. We can also combine 'Electricity' and 'Natural Gas' as 'Green Fuel'.

```
In [154]: for ft in ve_df.fuelType1.unique():  
           selection = ve_df['fuelType1'] == ft  
           if ft == 'Midgrade Gasoline':  
               ve_df.loc[selection, 'fuelType1'] = 'Regular Gasoline'  
           elif ft in ['Electricity', 'Natural Gas']:  
               ve_df.loc[selection, 'fuelType1'] = 'Green Fuel'  
           ve_df.fuelType1.value_counts()
```




```
Out[154]: Regular Gasoline    15732  
Premium Gasoline    12175  
Diesel              365  
Green Fuel          284  
Name: fuelType1, dtype: int64
```



# Correlation & Distribution Analysis




# Correlation & Distribution Analysis

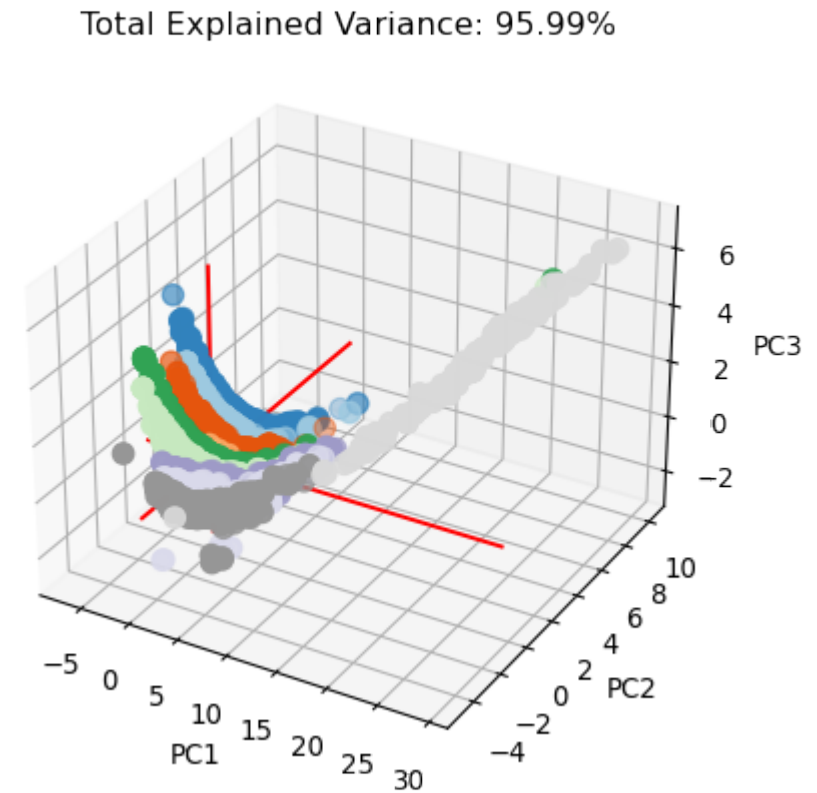


# Principal Component Analysis

---



- 3-component solution explains 96% variance in the original data
- For predictive modeling, we thought of avoiding PCA





1. Decision process: Classification




2. Error function: evaluate a model




3. Model optimization: adjusting model hyperparameters to get the best model

# Predictive Modeling




# Machine Learning




Ensemble Algorithms, Neural Network, SVM, Multinomial Logistic Regression

**Simple Neural Network**



**Deep Learning Neural Network**



● Input Layer    ● Hidden Layer    ● Output Layer




Model Evaluation and Comparison



Scalable ML using Dask

Source: <https://www.ibm.com/cloud/learn/machine-learning>



# Bias-Variance Tradeoff



Simultaneously minimize these two sources of error that prevent supervised learning algorithms from generalizing beyond their training set



High bias: underfitting  
High variance: overfitting



Regularization to control under or over fitting

Image source: [https://en.wikipedia.org/wiki/Bias%E2%80%93variance\\_tradeoff](https://en.wikipedia.org/wiki/Bias%E2%80%93variance_tradeoff)

# Model Comparison

---

1. **Ensemble Learning:** Gradient Boosting Trees (GBT), Random Forests (RF), Boosted Random Forests (RF), Extremely Randomized Trees (ETR)
2. **Neural Network:** Multi-layer Perceptron Classifier (MLP)
3. **Support Vector Classifier (SVC)**
4. **Multinomial Logistic Regression (MLR)**

**Note:** 70% train, 30% test with stratified split, Train and Test predictors scaled using MinMaxScaler for all the models

# Model Comparison: MLR

Out[219]:

	Data	F1 Score	Precision	Recall	Balanced Accuracy	Kappa Score
0	Train	0.331403	0.329090	0.509821	0.509821	0.189551
1	Validation	0.326589	0.326447	0.469016	0.483657	0.183054
2	Test	0.346034	0.353011	0.454485	0.409037	0.161838

In [220]: `print(report)`

	precision	recall	f1-score	support
1	0.13	0.31	0.18	584
2	0.25	0.61	0.35	354
3	0.32	0.53	0.40	566
5	0.49	0.19	0.27	3418
6	0.51	0.19	0.28	2506
7	0.32	0.35	0.33	527
8	0.06	0.49	0.11	164
9	0.14	0.42	0.21	370
10	0.96	1.00	0.98	76
micro avg	0.27	0.27	0.27	8565
macro avg	0.35	0.45	0.35	8565
weighted avg	0.42	0.27	0.28	8565

In [221]: `print(search_mlr.best_params_)`

```
{'C': 1, 'max_iter': 300, 'solver': 'saga', 'tol': 0.0001}
```



# Model Comparison: SVC

---

Out[20]:

	Data	F1 Score	Precision	Recall	Balanced Accuracy	Kappa Score
0	Train	0.370993	0.365296	0.577352	0.577352	0.231768
1	Validation	0.356233	0.355257	0.511651	0.532175	0.220769
2	Test	0.399593	0.386629	0.578398	0.578398	0.250525

In [21]: `print(report)`

	precision	recall	f1-score	support
1	0.46	0.46	0.46	584
2	0.24	0.75	0.36	354
3	0.41	0.71	0.52	566
4	0.13	1.00	0.24	2
5	0.63	0.23	0.34	3418
6	0.50	0.30	0.38	2506
7	0.33	0.54	0.41	527
8	0.06	0.37	0.10	164
9	0.14	0.43	0.21	370
10	0.97	1.00	0.99	76
accuracy			0.36	8567
macro avg	0.39	0.58	0.40	8567
weighted avg	0.50	0.36	0.37	8567

In [23]: `print(search_svm.best_params_)`

```
{'C': 1, 'break_ties': True, 'gamma': 0.8, 'tol': 0.01}
```

# Model Comparison: RF

---

Out[236]:

	Data	F1 Score	Precision	Recall	Balanced Accuracy	Kappa Score
0	Train	0.578722	0.555383	0.714078	0.714078	0.422408
1	Validation	0.436231	0.426816	0.533302	0.547823	0.263565
2	Test	0.408527	0.398950	0.420700	0.420700	0.227434

In [237]: `print(report)`

	precision	recall	f1-score	support
1	0.57	0.61	0.59	584
2	0.40	0.45	0.43	354
3	0.66	0.69	0.67	566
4	0.00	0.00	0.00	2
5	0.47	0.43	0.45	3418
6	0.35	0.35	0.35	2506
7	0.37	0.39	0.38	527
8	0.12	0.19	0.14	164
9	0.08	0.09	0.08	370
10	0.97	1.00	0.99	76
accuracy			0.42	8567
macro avg	0.40	0.42	0.41	8567
weighted avg	0.43	0.42	0.42	8567

In [238]: `print(search_rf.best_params_)`

```
{'class_weight': 'balanced_subsample', 'criterion': 'entropy', 'max_depth': None, 'max_features': 10, 'max_samples': 0.5, 'min_samples_leaf': 1, 'n_estimators': 500}
```

# Model Comparison: GBT

Out[22]:

	Data	F1 Score	Precision	Recall	Balanced Accuracy	Kappa Score
0	Train	0.702728	0.644762	0.839862	0.839862	0.540943
1	Validation	0.481861	0.452456	0.556968	0.568588	0.324939
2	Test	0.514768	0.471515	0.633385	0.633385	0.329062

In [23]: `print(report)`

	precision	recall	f1-score	support
1	0.50	0.66	0.57	584
2	0.36	0.67	0.47	354
3	0.55	0.81	0.65	566
4	0.50	1.00	0.67	2
5	0.64	0.30	0.41	3418
6	0.49	0.47	0.48	2506
7	0.39	0.54	0.45	527
8	0.15	0.48	0.23	164
9	0.15	0.41	0.22	370
10	0.97	1.00	0.99	76
accuracy			0.45	8567
macro avg	0.47	0.63	0.51	8567
weighted avg	0.53	0.45	0.46	8567

In [24]: `print(search_dask_lgbm.best_params_)`

```
{'boosting_type': 'gbdt', 'colsample_bynode': 1, 'colsample_bytree': 1, 'learning_rate': 0.05, 'max_bin': 127, 'max_depth': 10, 'min_child_samples': 20, 'min_data_in_bin': 3, 'n_estimators': 300, 'num_leaves': 31, 'path_smooth': 0.2, 'subsample': 1}
```

# Model Comparison: BRF

---

Out[20]:

	Data	F1 Score	Precision	Recall	Balanced Accuracy	Kappa Score
0	Train	0.457200	0.431045	0.675145	0.675145	0.331778
1	Validation	0.424028	0.404417	0.578920	0.602269	0.296177
2	Test	0.449043	0.417397	0.644578	0.644578	0.309786

In [21]: `print(report)`

	precision	recall	f1-score	support
1	0.46	0.62	0.53	584
2	0.33	0.72	0.45	354
3	0.49	0.80	0.60	566
4	0.10	1.00	0.17	2
5	0.64	0.26	0.37	3418
6	0.49	0.43	0.46	2506
7	0.38	0.50	0.43	527
8	0.14	0.66	0.23	164
9	0.18	0.46	0.26	370
10	0.97	1.00	0.99	76
accuracy			0.43	8567
macro avg	0.42	0.64	0.45	8567
weighted avg	0.52	0.43	0.43	8567

In [22]: `print(search_brf.best_params_)`

```
{'colsample_bynode': 1, 'colsample_bytree': 1, 'learning_rate': 0.05, 'max_bin': 255, 'max_depth': -1, 'min_child_samples': 20, 'min_data_in_bin': 3, 'n_estimators': 500, 'num_leaves': 63, 'path_smooth': 0, 'subsample': 0.9}
```

# Model Comparison: ETR

---

Out[33]:

	Data	F1 Score	Precision	Recall	Balanced Accuracy	Kappa Score
0	Train	0.533500	0.514130	0.709851	0.709851	0.401488
1	Validation	0.400826	0.395748	0.519259	0.537028	0.239507
2	Test	0.462869	0.428880	0.533174	0.533174	0.221032

In [34]: `print(report)`

	precision	recall	f1-score	support
1	0.55	0.61	0.58	584
2	0.39	0.53	0.45	354
3	0.62	0.67	0.64	566
4	0.40	1.00	0.57	2
5	0.46	0.36	0.40	3418
6	0.36	0.36	0.36	2506
7	0.34	0.40	0.37	527
8	0.12	0.24	0.16	164
9	0.10	0.16	0.12	370
10	0.96	1.00	0.98	76
accuracy			0.40	8567
macro avg	0.43	0.53	0.46	8567
weighted avg	0.42	0.40	0.40	8567

In [35]: `print(search_et.best_params_)`

```
{'class_weight': 'balanced', 'criterion': 'entropy', 'max_depth': None, 'max_features': 10, 'max_samples': 0.8, 'min_samples_leaf': 1, 'n_estimators': 300}
```

# Model Comparison: MLP

---

Out[21]:

	Data	F1 Score	Precision	Recall	Balanced Accuracy	Kappa Score
0	Train	0.526306	0.590153	0.513441	0.513441	0.385971
1	Validation	0.450087	0.489226	0.445315	0.453457	0.324013
2	Test	0.483473	0.530263	0.476535	0.428881	0.326520

In [22]: `print(report)`

	precision	recall	f1-score	support
1	0.57	0.59	0.58	584
2	0.49	0.38	0.43	354
3	0.57	0.72	0.64	566
5	0.52	0.67	0.59	3418
6	0.48	0.38	0.43	2506
7	0.53	0.37	0.43	527
8	0.39	0.09	0.15	164
9	0.25	0.09	0.13	370
10	0.97	1.00	0.99	76
micro avg	0.52	0.52	0.52	8565
macro avg	0.53	0.48	0.48	8565
weighted avg	0.51	0.52	0.50	8565

In [23]: `print(search_mlp.best_params_)`

```
{'hidden_layer_sizes': (512, 256), 'learning_rate_init': 0.001, 'max_iter': 100}
```

# Final Model

## Results: GBT

**ROC\_AUC** 0.622634

**F1 Score** 0.514768


**Precision** 0.471515

**Recall** 0.633385

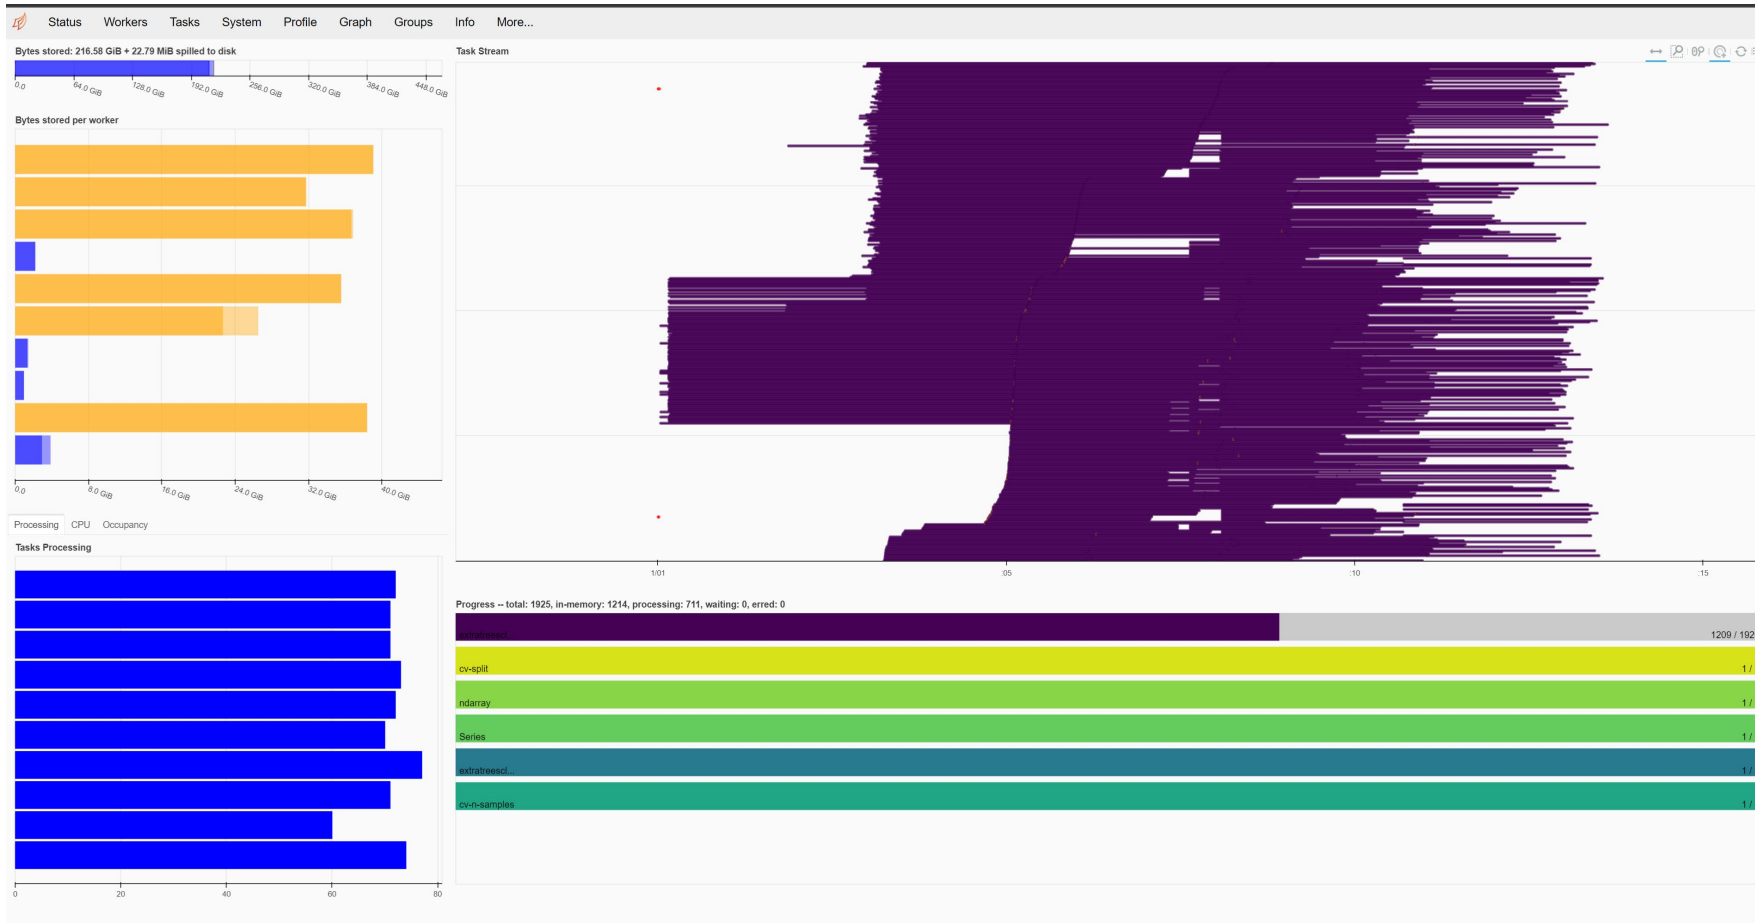
**Balanced Accuracy** 0.633385

**Kappa Score** 0.329062

Test Data



# Dask + Foundry



```
param_dict = {
    'n_estimators': [300, 500],
    'max_features': [3, 5, 7, 10],
    'max_depth': [6, 10, None],
    'max_samples': [0.8, 0.5, 0.9, None],
    'min_samples_leaf': [1],
    'criterion': ['gini', 'entropy'],
    'class_weight': ['balanced', 'balanced_subsample']
}

cv = RepeatedKFold(n_splits=5, n_repeats=1, random_state=random_state)
model = ExtraTreesClassifier(
    class_weight='balanced',
    random_state=random_state,
    n_jobs=-1
)

model_grid = DaskGCV(
    estimator=model,
    param_grid=param_dict,
    scoring=scoring_metrics,
    n_jobs=-1, cv=cv, refit=scoring_metrics['f1_'] + score_avg,
    return_train_score=True
)
```

## ETR Training:


- 216 GB memory!
- 40 cores
- 10 nodes
- ~10 mins of GridSearch



# Revisiting Research Questions

---


- Which predictors are most suited for predicting the EPA emission score?



# Revisiting Research Questions

---

- Which machine learning model works best for this data set? **Gradient Boosting Trees**
- Which vehicle class has poor emission ratings? **Pickup Trucks and Vans**



# Conclusions

- We were able to successfully answer the research questions
- 7 classification models are compared
- Correlation analysis and PCA-based dimensionality reduction were not helpful (possibly due to strong non-linearity)
- Additional data are required to improve the model results and address class imbalance.
- With Dask and Missouri S&T Foundry High-Performance Computing, we could save hours/days of computational time : Scalable ML is the way to go!
- This kind of modeling and analysis, if pushed to a production environment, could enable policymakers to address critical issues in the sustainable transportation industry.

# Thank you!

---

